

Deep Reinforcement Learning Exercise 02

November 17, 2024

Group:

- Marvin Kohnen
- Moritz Gehring
- Julius Ferber

1 Exercise 3.1

1.1 3.1a

Figure 1 shows the State Value functions and final optimal policy for the implemented GPI algorithm. Note that we reach state values of up to 45 in the best states (yellow).




Figure 1: State Value functions and final optimal policy visualization

1.2 3.1b

While using general policy iteration, the policy converges very stably and quickly in only a couple of iterations (see Figure 1). On the other hand, when using value iteration, the policy converges gradually and therefore much more slowly, but also equally reliably (see middle image of Figures 2 to 5). Of course, value iteration is much more computationally inexpensive than policy iteration, since

it only evaluates the value function once each iteration. As we can see though, the values reached in the final iteration are not as good as the ones reached using general policy iteration. Compared to the aforementioned state values for GPI (45 for the yellow states), the value iteration algorithm only reaches state values around 23. This is the case because the policy is updated right after on evaluation step, and not after the value function has converged.

In the end, both methods get to the same optimal policy (compare left most picture in Figures 1 and 5).




Figure 2: State Value functions and final optimal policy visualization after iteration 1




Figure 3: State Value functions and final optimal policy visualization after iteration 10




Figure 4: State Value functions and final optimal policy visualization after iteration 30




Figure 5: State Value functions and final optimal policy visualization after iteration 87

1.3 3.1c

Conceptual difference for GPI with Value function:

When using the action value function, we are taking the max over all actions for each state, and then updating the value function for each state. Regarding the State-value function, we're updating the value function for each state based on the expected return of each action in that state.

Moving on to the analysis:

Interestingly, the policy doesn't converge to the same optimal policy as in the previous exercises, but to a slightly different one. It looks like the policy using the action values has a smaller "scope", since the decision to go into the B cycle is made using action value functions for the whole column 3, while using state value functions, the A cycle is preferred for position (1, 3). The convergence seems to be of similar speed to the one using state value functions, since the number of policy iterations is again only 3, while the number of iterations to evaluate the action values is 89 (see 6), like when using state value functions.

```
31 iterations to evaluate action values
89 iterations to evaluate action values
89 iterations to evaluate action values
89 iterations to evaluate action values
```

Figure 6: Logging output monitoring iterations to evaluate action values.

Optimal Policy Visualization (Iteration: 3)




Figure 7: Optimal policy visualization




Figure 8: Action values by direction




Figure 9: Optimal policy visualization

2 Exercise 3.2

2.1 3.2.a

The initial policy is defined with the following probabilities:

- Moving RIGHT: 35%
- Moving DOWN: 35%
- Moving LEFT: 15%
- Moving UP: 15%

The rationale behind this policy is that, as the goal is located at the bottom right of the lake and the starting point is at the top left, the general strategy should be to move towards the bottom-right direction.

When applying Monte Carlo prediction in the context of Reinforcement Learning, the characteristics that an initial policy must have are:

1. **Coverage of entire State-Action Space:** The policy should be able to visit all relevant states and state-action pairs. Ensuring that all areas of the state-action space are covered ensures that the estimated values are accurate.

2. **Consistency with Problem Constraints:** The initial policy should respect any constraints inherent in the problem domain (eg. using only valid movement commands in this case).

3. **Exploratory nature** The initial policy should ensure sufficient exploration of the state space.

Our policy adequately reflects all 3 characteristics.

Moving on to the Analysis part:

A state value is only the average action value in that state across all actions and the policy. However, if the state has an action that guarantees a reward (e.g. the action to the left of the goal, in this case with a value of 0.6 - see Figure 10), that fact is not accurately represented by just the state value function. Similarly, states far from the goal are all evaluated very close to each other in value (in this case 0), even though some of them are clearly superior (e.g. tile 6 compared to tile 4), but since the agent generally fails to retrieve a reward very often, they have almost the same value.

2.2 3.2.b

The visualization of the On-Policy Monte Carlo can be seen in Figure 11 and the optimal policy in Figure 12. We also plotted the heatmap for each action in Figure 13

Even though in (3,2), the action RIGHT should give a reward of 1 every time, because it is slippery and there is a randomness factor, its state-action value is less than 1 (Figure 13). Controversely, the action value for the action DOWN is actually sometimes greater than the one for RIGHT. This is because




Figure 10: Monte Carlo Heatmap

the discounting factor ‘gamma’ is 1 (i.e. there is no discounting), and because when using the action DOWN in (3,2), there is a chance that the agent slips into the goal state, and if not, it either slips into another safe state (there are no adjacent holes) or stays in (3,2). Therefore, the calculated action values can only be relied upon to a certain extent, as the best action in (3,2) is obviously to move RIGHT, but due to the non-deterministic environment dynamics, this will not always be represented in the state-action values.

2.3 3.3

I’m so sorry boys, but i’m **fucking cooked**.

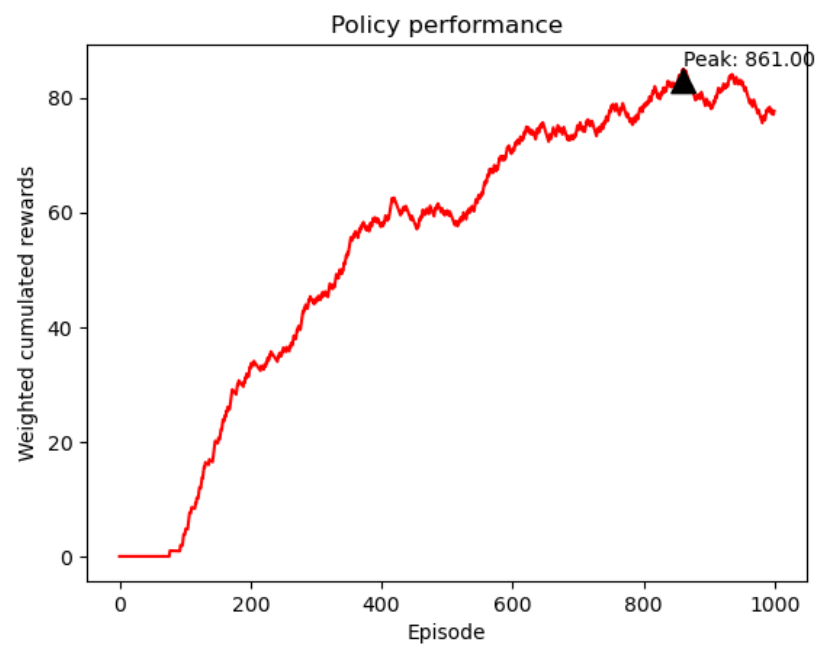


Figure 11: Monte Carlo Policy Performance

Optimal Policy

←	↑	↑	↑
←	←	↓	←
↑	↓	↓	←
←	↓	↓	←

Figure 12: Optimal Policy for On-policy Monte Carlo Control



Figure 13: Heatmap for each Action